

Ground Rules: Do all problems below. Solve them either by yourself or in teams. Typeset your answers and hand in a PDF file electronically to Canvas. You can look up things on the Internet; refrain from copying solutions straight-up.

Problem 1. *LSH for Euclidean Distance.* We'll further study an LSH family for Euclidean distance in \mathbb{R}^d .

- (i) To warm up, consider the following problem: Fix $w > 0$ and draw $s \in [0, w]$ uniformly at random. Prove that if we define $f(x) = \lfloor \frac{1}{w}(x + s) \rfloor$, then

$$\Pr[f(x) = f(y)] = \max\left(0, 1 - \frac{1}{w}|x - y|\right).$$

- (ii) We define a hash family \mathcal{H}_r for \mathbb{R}^d by describing a process to generate a (random) hash function from this family: To draw a random hash function h from \mathcal{H}_r , draw a random number $s \in [0, r]$ uniformly at random and form a vector $\mathbf{v} = \langle v_1, v_2, \dots, v_d \rangle$ such that $v_i \sim N(0, 1)$. Then, $h: \mathbb{R}^d \rightarrow \mathbb{Z}$ is given by

$$h(\mathbf{x}) = \left\lfloor \frac{1}{r}(\mathbf{x}^\top \mathbf{v} + s) \right\rfloor$$

Show that for fixed r and c , the hash family \mathcal{H}_r is (r, cr, p_1, p_2) -locality-sensitive by deriving p_1 and p_2 . It is okay to leave the expressions for p_1 and p_2 in terms of definite integrals involving the Gaussian pdf; however, you'll want to show that $p_1 > p_2$.

Problem 2. *Dual Binary Search and Better Merge Sort.* Consider the implementation of $\text{kth}(A, B, k)$ in the given handout. We will analyze its cost and use it to build merge sort with better span.

- (i) Show that $\text{kth}(A, B, k)$ runs in $O(\log |A| + \log |B|)$ work and span. (*Hint:* Can we guarantee that something goes down in half in each recursive call?)
- (ii) We'll start by improving the span bound of merge. To merge two sorted sequences, we'll split them into \sqrt{n} equal-sized pieces (previously, we split them into 2 equal-sized pieces). Describe how you would use the kth routine to accomplish this. Give a (pseudocode) implementation of the upgraded merge.
- (iii) What is the work recurrence for our upgraded merge? How about span? (*Hint:* You should get $W(n) = \sqrt{n}W(\sqrt{n}) + \dots$ and $S(n) = S(\sqrt{n}) + \dots$)
- (iv) Let $f(n)$ be $o(n)$ —this reads little- O of n . Solve the following recurrence: $W(n) = \sqrt{n}W(\sqrt{n}) + f(n)$. (*Hint:* $\Theta(n)$)
- (v) Now that we know how to solve the recurrences, what are the work and span bounds for our upgraded merge?
- (vi) If merge sort now uses the improved merge routine, what is its overall work and span?

Problem 3. *Quick Sort's Span.* In class, we made quick sort parallel by showing that both the partitioning and concatenation steps can be accomplished in $O(n)$ work and $O(\log n)$ span. Show that the span of this version of quick sort is $O(\log^2 n)$ whp. Notice that the work bound follows from the argument we did in the previous assignment.

Problem 4. String Comparison. Derive a parallel algorithm $\text{compare}(X, Y)$ that lexicographically compares two given strings $X = x_1x_2 \dots x_n$ and $Y = y_1y_2 \dots y_m$ and returns the following result

$$\text{compare}(X, Y) = \begin{cases} -1 & \text{if } X < Y \\ 0 & \text{if } X = Y \\ +1 & \text{if } X > Y \end{cases}$$

Your algorithm should run in $O(\min(n, m))$ work and $O(\log \min(n, m))$ span.

Problem 5. Parallel Closest Pair. In the sequential setting, the closest pair of points in a set of n 2-d points can be found in $O(n \log n)$ time using divide and conquer. Make this algorithm parallel so that it runs in $O(n \log n)$ work and at most $O(\log^3 n)$ span. (*Hint:* The points on the boundaries of the two sides of a recursive call can be examined in parallel. Also, we can prove that there is no point in looking at more than a constant number of points.)