
Contemporary Algorithms T.2/2019-20

Basic Information/Logistics

- Website: <https://cs.muic.mahidol.ac.th/courses/calgo>
- This course: highlights of useful data structures & algorithmic ideas from the past 50 years.
- Schedule will be made as we go along.
- No lecture notes. You'll scribe!

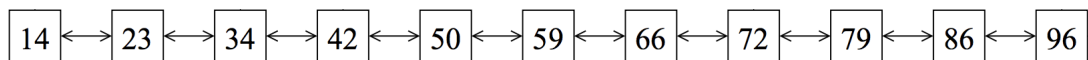
Week 1: Ordered Maps

- Think TreeMap in Java
- The keys are ordered supporting add/remove/update/lookup
- Goal: for most operations, take $O(\lg n)$ time or faster

The themes: approximating a perfectly balanced structure

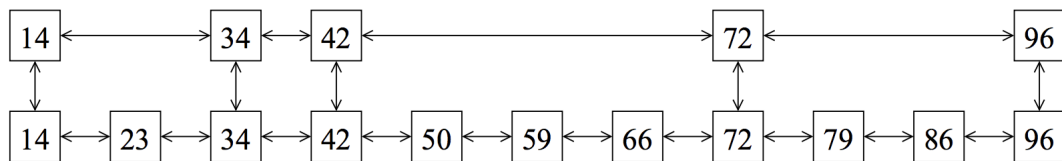
Skip Lists

Starting simple,



Sorted Linked List: $O(n)$

Can we do better?



2-level Linked List: $m + \frac{n}{m} \leq \sqrt{n} + \frac{n}{\sqrt{n}} = 2\sqrt{n}$

3-LL: $3\sqrt[3]{n}$

k-LL: $k\sqrt[k]{n}$

Pick k so that search time is minimized, we can get search time = $2\lg n$ and becomes a perfect binary tree.

Maintaining this exactly is too rigid, we relax by flipping a coin.

For each item, we toss a fair coin. If the coin turns head, we promote that item up. By this, we also get approximately $\lg n$ layers.

To insert: find right spot, promote the new element (by 0.5 prob) and cut up necessary siblings to remain the structure.

$\mathbb{E}[\text{search cost}] = ?$ Think backwards from bottom to top.

For one layer, $\mathbb{E}[\text{walk across}] = \mathbb{E}[\text{cannot walk up}] = \mathbb{E}[\# \text{ toss tails until heads}] = \frac{1}{p} = 2$

We expect $\lg n$ layers, so search cost = $2\lg n$.

Definition: An even E_α occurs with **high probability (whp)** if for any α , $P[E_\alpha] \geq 1 - \frac{c_\alpha}{n^\alpha}$ where c_α is a constant that depends only on α .

Thm: Every search costs $\Theta(\lg n)$ whp.

Proof ideas:

Analyze search backwards from bottom to top layers.

- Start at the found element at the bottommost layer.
- If we get tail, walk across to the left. If we get head, walk up the tree.
- Stop when reach the topmost layer.

Need to show two things

- (i) Walk up is $O(\lg n)$.
- (ii) Walk across is $O(\lg n)$.

Lemma (i): Skip list has $O(\lg n)$ levels whp. (Showing walk up is $O(\lg n)$)

Proof:

$$Pr[\text{a key } k_i \text{ grows taller than } 100 \lg n] = \frac{1}{2^{100 \lg n}} = \frac{1}{n^{100}}$$

$$\begin{aligned} Pr[\text{a skiplist has } \geq 100 \lg n \text{ levels}] &= Pr[\exists i, k_i \text{ grows to height } \geq 100 \lg n] \\ &\leq Pr[k_1 \text{ is too tall}] + Pr[k_2 \text{ is too tall}] + \dots \\ &= n \times \frac{1}{n^{100}} \\ &= \frac{1}{n^{99}} \end{aligned}$$

Therefore,

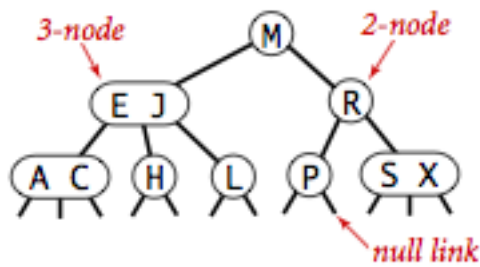
$$\begin{aligned} Pr[\text{a skiplist has } < 100 \lg n \text{ levels}] &= 1 - Pr[\text{a skiplist has } \geq 100 \lg n \text{ levels}] \\ &\geq 1 - \frac{1}{n^{99}} \end{aligned}$$

Lemma(ii) : Out of $1000 \lg n$ flips, you get $\geq 100 \lg n$ heads up. (Showing walk across is $\Theta(\lg n)$)

Proof:

$$\begin{aligned} Pr[\text{getting exactly } 100 \lg n \text{ heads}] &= \binom{1000 \lg n}{100 \lg n} \underbrace{\frac{1}{2}^{900 \lg n}}_{\text{Tail}} \underbrace{\frac{1}{2}^{100 \lg n}}_{\text{Head}} \\ Pr[\text{getting } < 100 \lg n \text{ heads}] &= \binom{1000 \lg n}{100 \lg n} \underbrace{\frac{1}{2}^{900 \lg n}}_{\text{Tail}} \\ &\leq \frac{e^{1000 \lg n}}{100 \lg n} \left(\frac{1}{2} \right)^{900 \lg n} \\ &= 2^{(\lg(10e))100 \lg n - 900 \lg n} \\ &= 2^{100 \lg(10e) - 900} \\ &= n^{-\alpha} \\ &= \frac{1}{n^\alpha} \end{aligned}$$

2-3 search trees



Anatomy of a 2-3 search tree

In **BST**, a node has 1 key and 2 links, we call this a **2-node**.

In **2-3 search tree**, we also allow a node with 2 keys and 3 links which we call a **3-node**.

Definition: A 2-3 search tree is either empty or

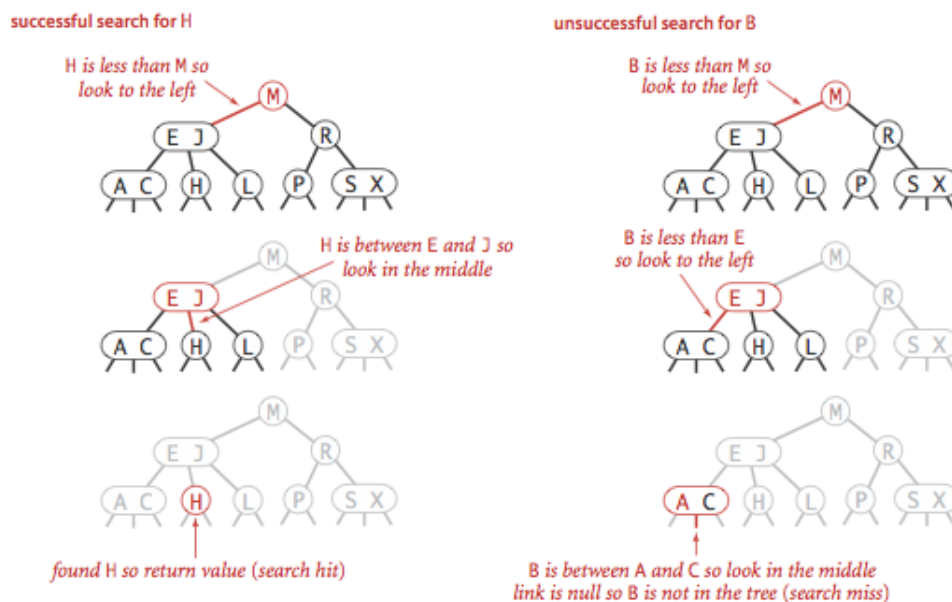
- A 2-node (with one key, two links: a left link with a 2-3 search tree with smaller keys, a right link with a 2-3 search tree with larger keys.)
- A 3-node (with two keys, three links: a left link with smaller keys, a middle link with keys in between, a right link with larger keys.)

note: a *null link* is a link to an empty tree.

Null links of a *perfectly balanced* 2-3 search tree are at the same distance from the root.

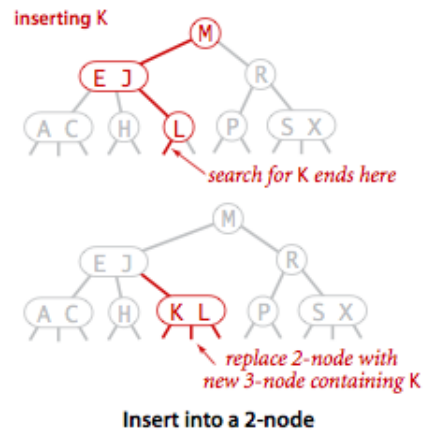
Goal: Any insert operation must still make a 2-3 search tree perfectly balance. (From now, we'll use the term *2-3 tree* to refer to a perfectly balanced 2-3 tree.)

Search: the search algorithm is the same as that of BSTs. We either get a search hit or miss.

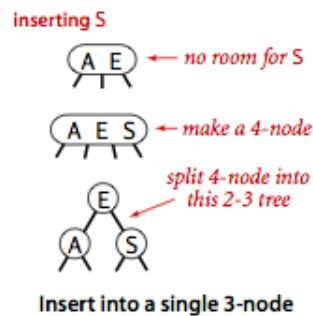


Search hit (left) and search miss (right) in a 2-3 tree

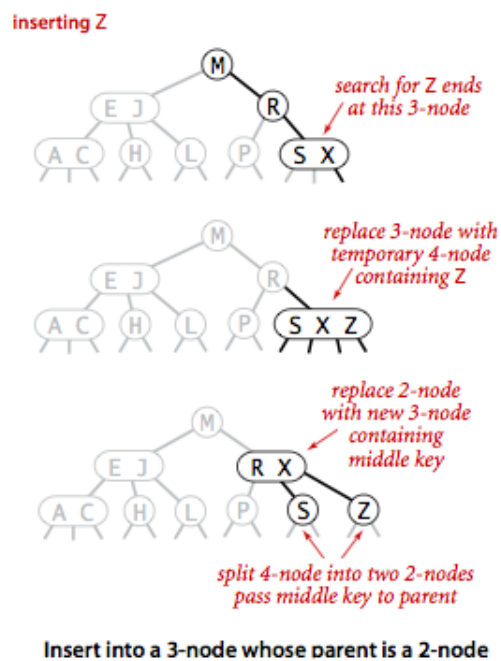
Insert into a 2-node: if we do an unsuccessful search and terminate at a 2-node at the bottom, just replace the node with a 3-node.



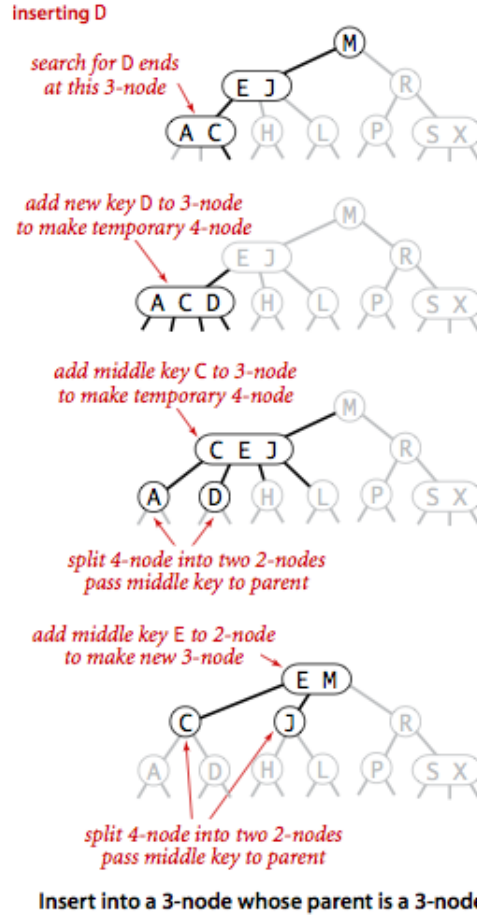
Insert into a tree of a single 3-node:



Insert into a 3-node whose parent is a 2-node:



Insert into a 3-node whose parent is a 3-node:



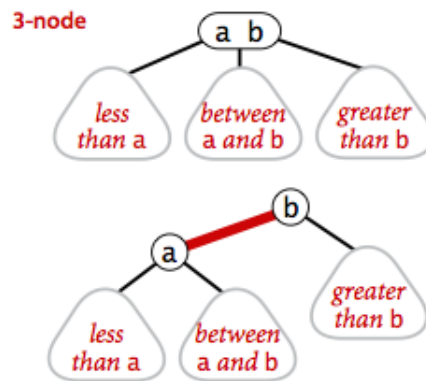
Proposition: Search and insert in a 2-3 tree with N keys are guaranteed to visit at most $\log(N)$ nodes.

Proof: The height of an N -node 2-3 tree is between $\lceil \log_3 N \rceil$ (if the tree is all 3-nodes) to $\lfloor \log_2 N \rfloor$ (if the tree is all 2-nodes).

Red-black BSTs

We'll implement a 2-3 tree using a *red-black BST* representation. We have two types of links: red and black.

- *Red* links bind two 2-nodes to represent 3-nodes. Red links **lean left**.
- *Black* links bind together the 2-3 tree.
[see a figure below.]



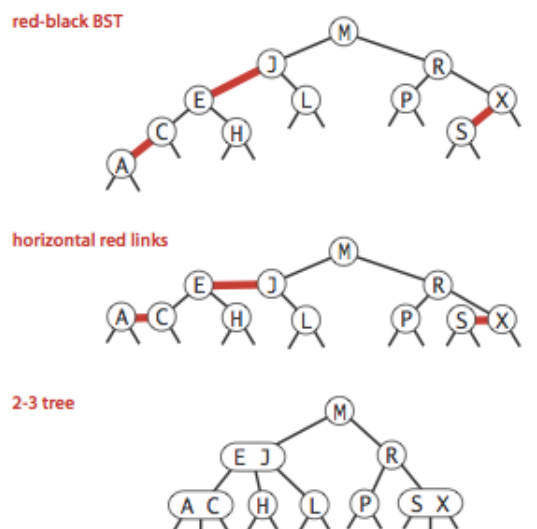
We can see that this representation above (the one with red link) allows us to use code from the standard BST search without modification.

Definition: red-black BSTs are BSTs that have red and black links that satisfy:

- Red links lean left.
- No node has 2 red links connected to it. (no consecutive red links)
- Every path from the root to the null link (a link to an empty tree) has the same number of black links. (**perfect black balance**)

1-1 correspondence between red-black BSTs and 2-3 trees:

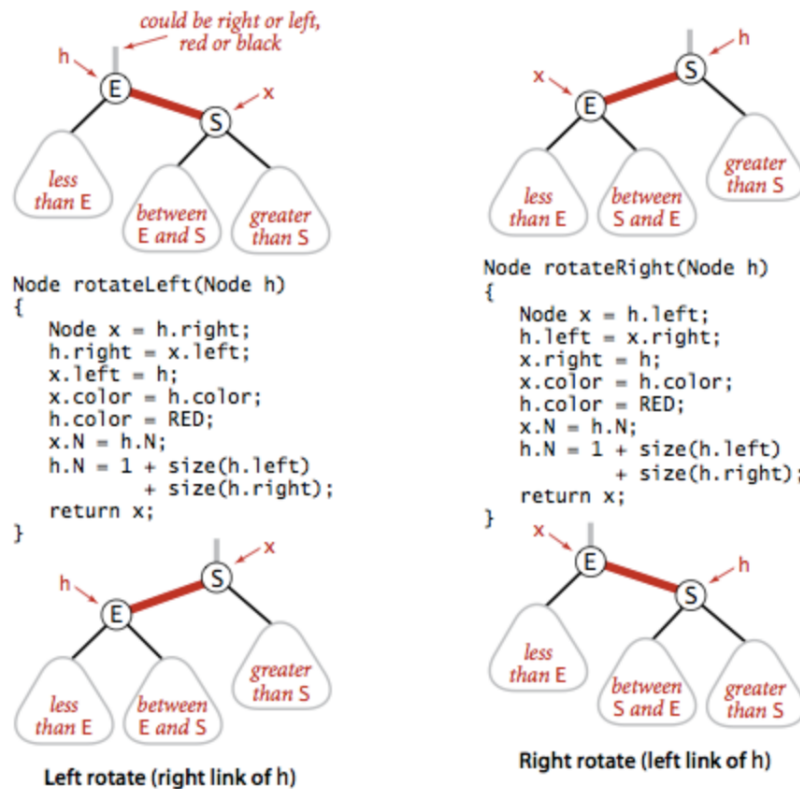
- If we draw the red links horizontally in a red-black BST, all the null links are the same distance from the root.
- If we collapse the nodes connected in red links, we get a 2-3 tree.



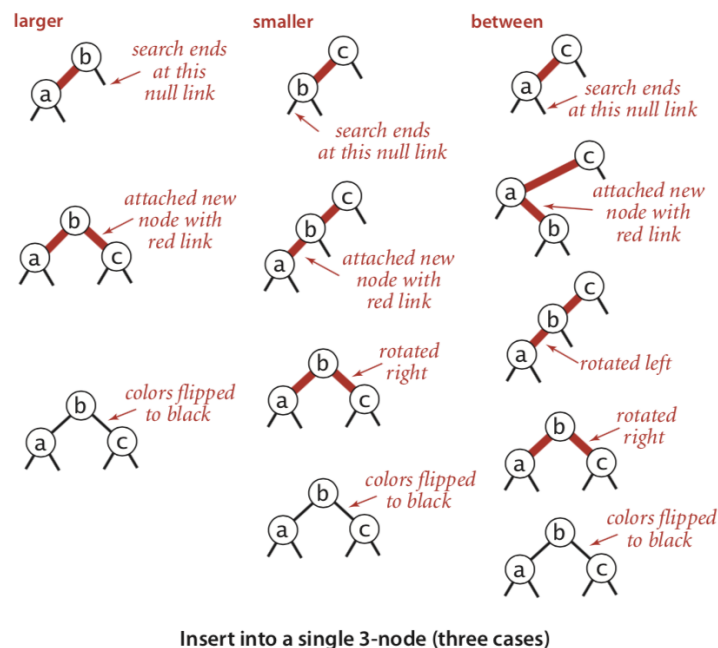
1-1 correspondence between red-black BSTs and 2-3 trees

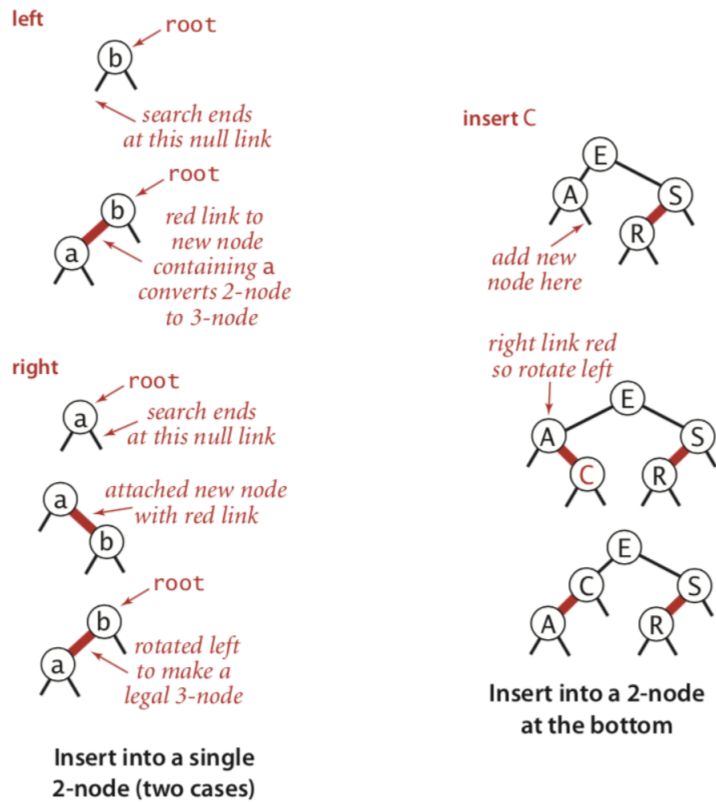
Rotations:

- We use rotations to maintain the 1-1 correspondence between red-black BSTs and 2-3 trees.



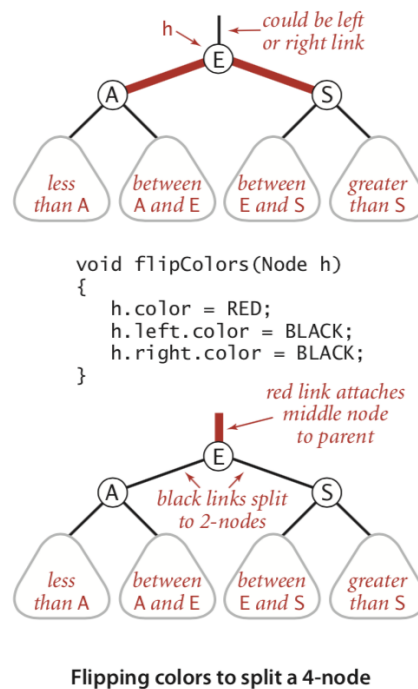
- We also use rotations to maintain the other 2 properties: no consecutive red links and red links must only lean left.





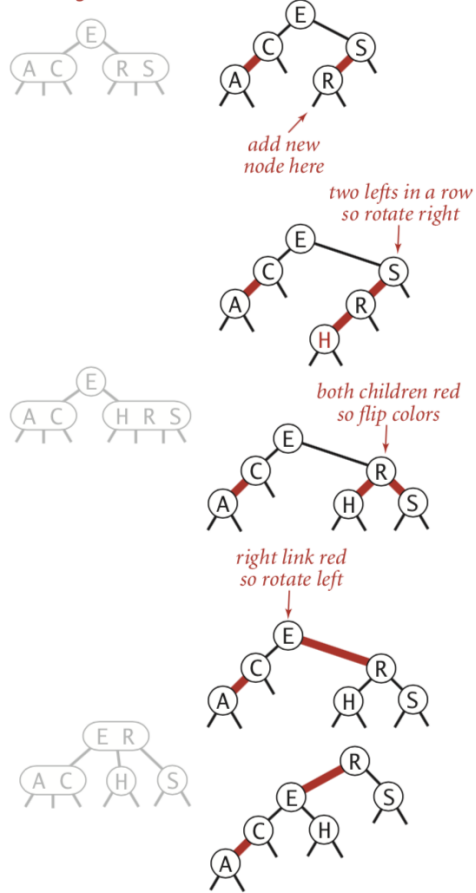
Flipping colors:

- we flip color to preserve the property **perfect black balance** in the red-black BSTs.



Insert into a 3-node at the bottom:

inserting H



Insert into a 3-node at the bottom