

Lecture 13: Gradient Descent and Stochastic Gradient Descent

DATE

Lecturer: Kanat Tangwongsan

Scribe: Kriangsak Thuiprakhon

1 Gradient Descent

1.1 Remarks

- (i) What the heck is $\eta \nabla f(x)$? The second-order Taylor's expansion :

$$f(y) \approx f(x) + f'(x) \cdot (\vec{y} - \vec{x}) + \frac{1}{2} f''(x) (\vec{y} - \vec{x})^2$$

In $d \geq 2$,

$$f(\vec{y}) \approx f(\vec{x}) + \nabla f(\vec{x})^T (\vec{y} - \vec{x}) + \frac{1}{2} (\vec{y} - \vec{x})^T \nabla^2 f(x) (\vec{y} - \vec{x})$$

where $\nabla^2 f(x)$ is Hessian $Hf(x)$. This potentially complex Hf can be approximated by an extremely simple term $\frac{1}{\eta} I$, where I is the identity matrix. That is,

$$\tilde{f}(\vec{y}) = f(\vec{x}) + \nabla f(\vec{x})^T (\vec{y} - \vec{x}) + \frac{1}{2\eta} \|\vec{y} - \vec{x}\|_2^2$$

This turns out to be convex for some complex reasons. The question now is how to minimize $\tilde{f}(\vec{y})$?. To do this, by convexity, set $\nabla_y \tilde{f}(\vec{y}) = 0$. Then,

$$\nabla_y \tilde{f}(\vec{y}) = \nabla f(\vec{x}) + \frac{1}{\eta} (\vec{y} - \vec{x}) = 0 \iff \vec{y} - \vec{x} = -\frac{1}{\eta} \nabla f(\vec{x})$$

This technique is called **minimization of local approximations** where $f + \nabla$ is local approximation and $\frac{1}{2\eta} \|\vec{y} - \vec{x}\|$ is proximity.

- (ii) In practice, how to choose η_t ? Fixing η_t throughout an approximation is not always a good idea. Therefore, a popular heuristic is to have a decaying η_t
- (iii) For further refinement, for instance, the function f needs to be *strongly* convex (i.e., the function must "bend").

1.2 Special Case Popular in Machine Learning/ Deep Learning

- minimize a loss function $f(\vec{\theta}) = \frac{1}{n} \sum_{i=1}^n f_i(\vec{\theta})$
- Example: In linear regression with l_2 loss. Data points $(\vec{x}_{i,j}, y_i)$ where y_i is a label. The loss function is given by,

$$\text{loss} = \sum_{i=1}^n (y_i - \vec{x}_i^T \vec{\theta})^2$$

Minimizing this loss function is equivalent to minimizing

$$f(\vec{\theta}) = \frac{1}{n} \sum_{i=1}^n (y_i - \vec{x}_i^T \vec{\theta})^2$$

- Notice that running GD takes $O(nd)$ per iteration. This is too costly, there shall be a way to reduce this.
- Fun fact: $\mathbb{E}_i[f_i(\vec{\theta})] = f_i(\vec{\theta})$ and $\nabla_{\theta} f_i(\vec{\theta}) = \frac{1}{n} \sum_{i=1}^n \nabla_{\theta} f_i(\vec{\theta})$ This fact gives rise to Stochastic Gradient Descent (SGD) whereby space requirement and time complexity per iteration is reduced.

2 Stochastic Gradient Descent

We can express the algorithm as the following.

Algorithm 1 Stochastic Gradient Descent Algorithm

```

function GRADIENTDESCENT( $f, x_0$ )
  for  $t = 1, 2, \dots, T - 1$  do
     $x_t \leftarrow x_{t-1} - \eta_t \nabla \mathbb{E} f_i(x_{t-1})$ 
  return  $\hat{x} = \frac{1}{T} \sum_t x_t$ 

```

SGD modifies this so that.

Algorithm 2 Stochastic Gradient Descent Algorithm

```

function GRADIENTDESCENT( $f, x_0$ )
  for  $t = 1, 2, \dots, T - 1$  do
    choose an  $i$  at random or round robin.
     $x_t \leftarrow x_{t-1} - \eta_t \nabla f_i(x_{t-1})$ 
  return  $\hat{x} = \frac{1}{T} \sum_t x_t$ 

```

Theorem 2.1. Let x^* be the minimiser of $f : \mathbb{R}^n \rightarrow \mathbb{R}$. If f is convex and differentiable and satisfies $\|\nabla f_i(x)\|_2 \leq G$ for all $x \in \mathbb{R}^n$, then setting $T = \frac{G^2}{\epsilon^2} \|x_0 - x^*\|_2^2$ and $\eta_t = \eta = \frac{\|x_0 - x^*\|_2}{G\sqrt{T}}$ gives $\mathbb{E}f(\hat{x}) \leq f(x^*) + \epsilon$.

This is the same old proof. The interesting bits are:

Claim 2.2. Let $\Phi_t = \frac{1}{2\eta} \|x_t - x^*\|_2^2$. Then, $\mathbb{E} \left[f(\vec{x}_t + (\Phi_{t+1} - \Phi_t)) \right] \leq f(x^*) + \frac{1}{2} \eta G^2$

Condition \mathbb{E} on the history until iteration that produced \vec{x}_t

$$\begin{aligned} \mathbb{E}[f(\vec{x}_t) + \Phi_{t+1} - \Phi_t] &\leq \mathbb{E} \left[f(\vec{x}_t) + \frac{1}{2} \eta \left(\underbrace{\|\vec{x}_{t+1} - \vec{x}_t\|_2^2}_{\Delta x} + 2\Delta x^T(\vec{x}_t - x^*) \right) \right] \\ &\leq \mathbb{E}[f(\vec{x}_t)] + \frac{1}{2} \eta G^2 - \nabla f(\vec{x})^T(\vec{x}_t - x^*) \\ &= \frac{1}{2} \eta G^2 + \underbrace{f(\vec{x}_t) + \nabla f(\vec{x})^T(x^* + \vec{x}_t)}_{\leq f(x^*)} \end{aligned}$$

Note:

$$\begin{aligned} \Delta x &= -\eta f_i(\vec{x}_t) \\ \Rightarrow \|\Delta x\|_2^2 &\leq \eta^2 G^2 \\ \Rightarrow \mathbb{E}_i[\Delta x] &= -\eta \mathbb{E}[\Delta_{f_i}(\vec{x}_t)] = -\eta \Delta f(\vec{x}_t) \end{aligned}$$

Other tricks used to implement SGD:

- $\mathbb{E}_i[f_i(\vec{x}) = \mathbf{f}_i(x)$ but potentially not concentrated. Therefore, This potentially does not converge as fast as performing the full gradient.
- To improve further, Pick a "batch" of indices. Say, $B \leq [n]$ and use an update rule such that

$$\vec{x}_t = \vec{x}_{t-1} - \frac{\eta}{|B|} \sum_{i \in B} \nabla f_i(\vec{x}_{t-1}).$$

$$\text{Note: } \mathbb{E} \left[\frac{1}{|B|} \cdot \sum_{i \in B} \nabla f_i(\vec{x}) \right] = \frac{1}{|B|} \sum_{i \in B} \nabla f_i(\vec{x}) = \nabla f(\vec{x})$$

By doing this, minibatch sampling, the variance is reduced by about $\frac{1}{|B|}$ and the cost per iteration shrinks down to $O(|B| \cdot d)$. This is considered a good compromise in practice since the convergence rate, (i.e., $f(\hat{x}) - f(x^*)$), is $O\left(\sqrt{\frac{|B|}{T} + \frac{|B|}{T}}\right)$ as opposed to $O\left(\frac{1}{T}\right)$ in a normal GD algorithm.